

Deep Learning

THE GOOD THE BAD and THE UGLY

Qiyang Hu

UCLA Office of Advanced Research Computing

November 8th, 2021

What makes our Deep Learning workshops special?

- Different from other learning resources
 - **Broad**: high-level descriptive review
 - **Intuitive**: avoiding math to explain the fundamentals
 - **Practical**: jupyter notebook examples on basic techniques
- Different from my previous LDL series
 - More **insights**: esp. from AI etc.
 - More **critical** perspectives
 - More research in **Science**
- Plans:
 - This quarter: General discussions on DL, Learning Mech, PyTorch
 - Next quarter: Specific topics on CNNs, GANs, RNN...

Neutral

General Intro



- What's ANN?
- DL, ML and AI

01

Good

DL as hopes



- Why DL is so powerful?
- Exciting achievements
- Mutual driving forces

02

Bad

DL as hypes



- Difficulties in application perspective
- Illusions in theoretical perspective

03

"Ugly"

DL like deep oceans



- Cracking the black box
- Thinking outside the box
- What about the future?

04

Neutral

General Intro



- What's ANN?
- DL, ML and AI

01

Good

DL as hopes



- Why DL is so powerful?
- Exciting achievements
- Mutual driving forces

02

Bad

DL as hypes



- Difficulties in application perspective
- Illusions in theoretical perspective

03

"Ugly"

DL like deep oceans




- Cracking the black box
- Thinking outside the box
- What about the future?

04




What is Neural Network?

- Recap for simple linear classification problem



$$z = w_1 x_1 + w_2 x_2 + b$$







$$y = \sigma(z) = \begin{cases} 1 & \text{if } z > 0 \\ 0 & \text{if } z < 0 \end{cases}$$




Artificial Neuron and Biological Neuron

**McCulloch-Pitts
(MCP) neuron model**






Neural Networks ~ piling/stacking logistic-regression classifiers



Deep neural networks


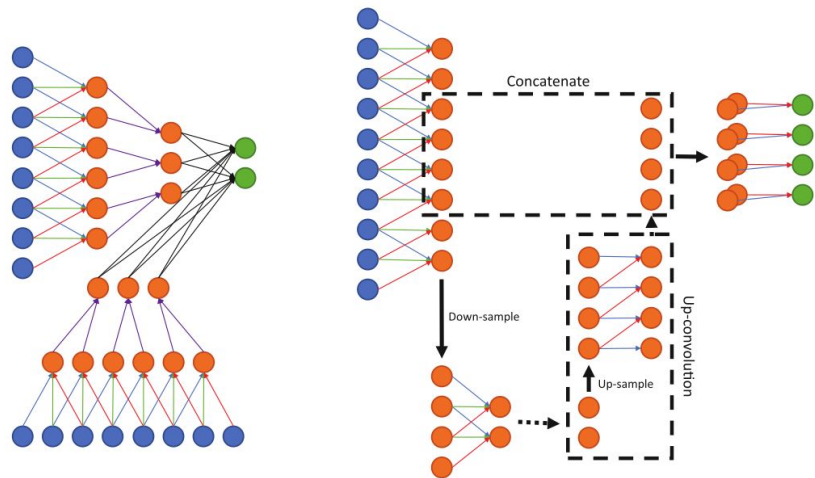



- LeNet-5 (1998)



Year	CNN	Developed by	Place	Top-5 error rate	No. of parameters
1998	LeNet(8)	Yann LeCun et al			60 thousand
2012	AlexNet(7)	Alex Krizhevsky, Geoffrey Hinton, Ilya Sutskever	1st	15.3%	60 million
2013	ZFNet()	Matthew Zeiler and Rob Fergus	1st	14.8%	
2014	GoogLeNet(19)	Google	1st	6.67%	4 million
2014	VGG Net(16)	Simonyan, Zisserman	2nd	7.3%	138 million
2015	ResNet(152)	Kaiming He	1st	3.6%	

Deep neural *networks*



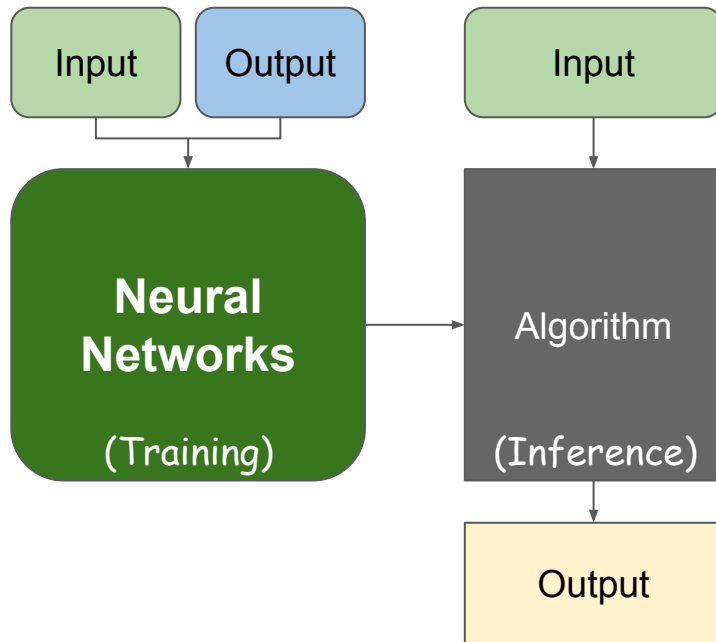


What is Deep Learning?

Traditional Programming



Deep Learning



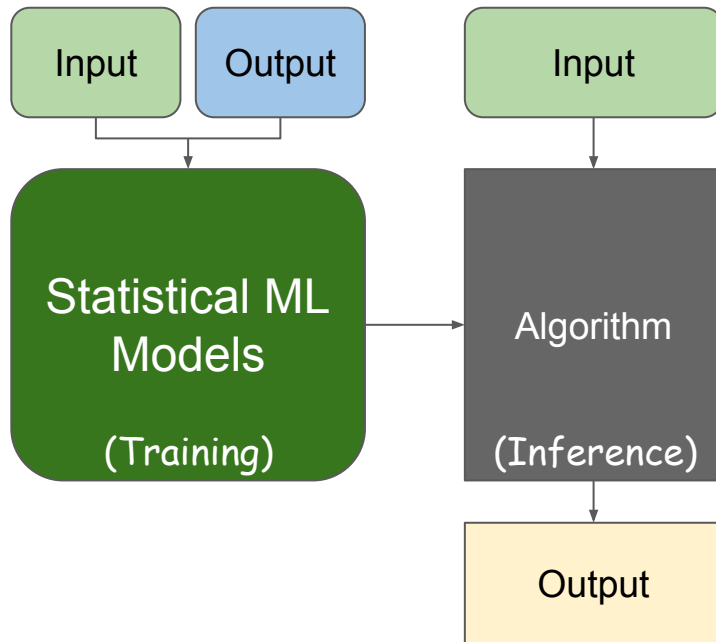


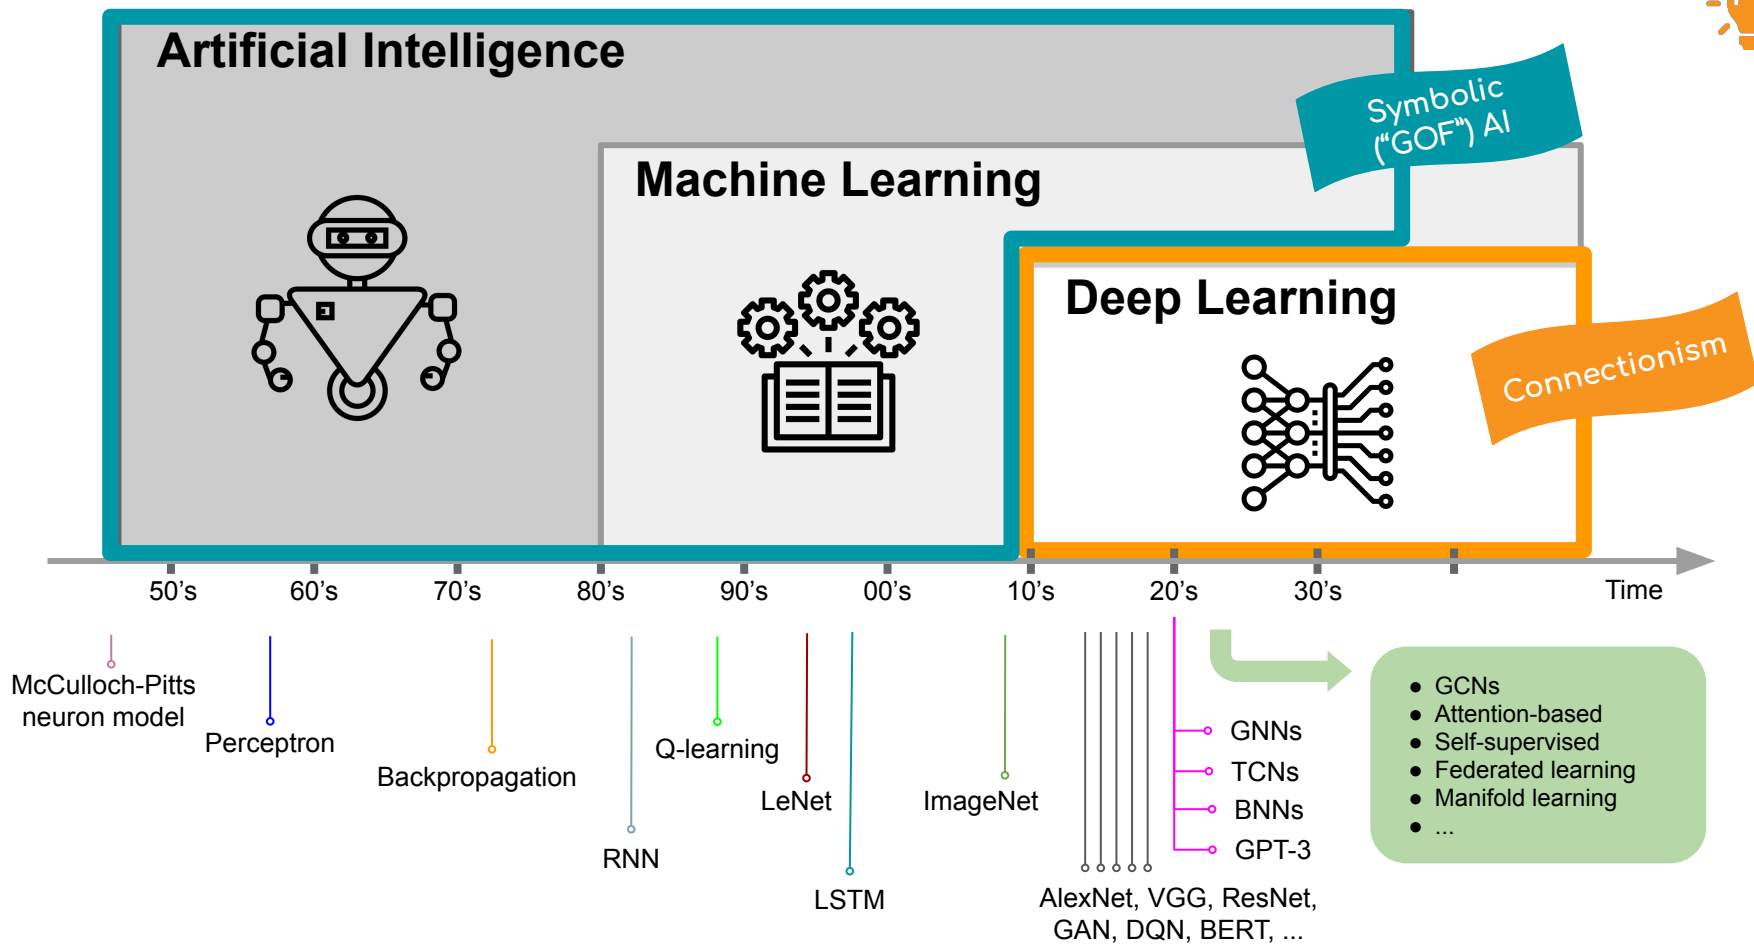
What is Machine Learning?

Traditional Programming



Machine Learning





Neutral

General Intro



- What's ANN?
- DL, ML and AI

01

Good

DL as hopes



- Why DL is so powerful?
- Exciting achievements
- Mutual driving forces

02

Bad

DL as hypes



- Difficulties in application perspective
- Illusions in theoretical perspective

03

"Ugly"

DL like deep oceans




- Cracking the black box
- Thinking outside the box
- What about the future?

04



DATA CENTER & INTERNET



SEARCH
RECOMMENDATION

ML DOMAIN



RESEARCH HIGHER EDUCATION



SIMULATION
LANGUAGE



HEALTHCARE



DRUG DISCOVERY
GENOMICS



FINANCE



QUANT FINANCE
FRAUD DETECTION



AUTOMOTIVE




IMAGE DETECTION,
CLASSIFICATION





A deeper view on connectionism

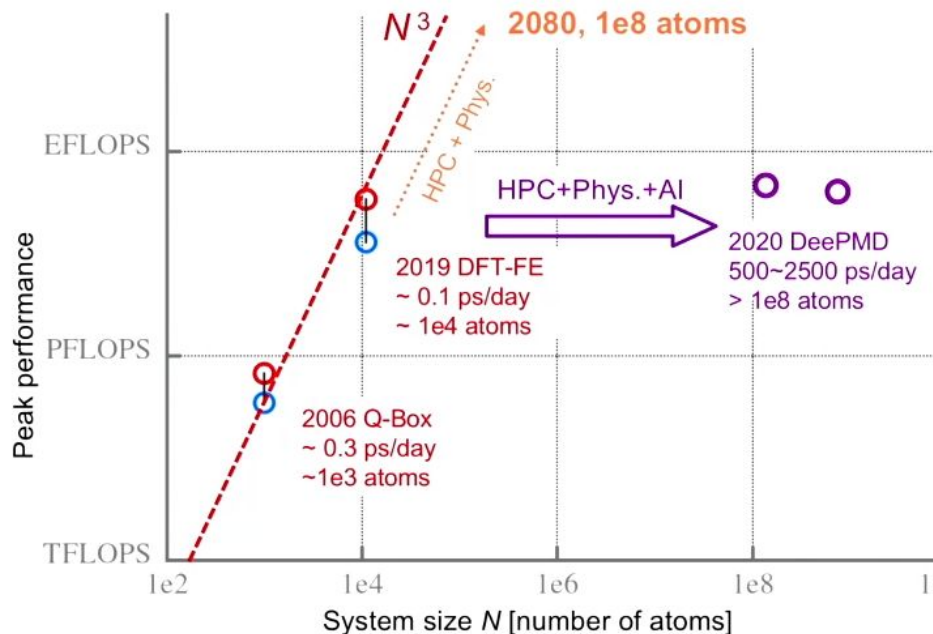
- Knowledge representation
 - Concepts: vectors of neural activity
 - Relationships between concepts: weight matrices
- Automatic generalization of learning
 - Facilitates analogical reasoning
 - Hierarchical knowledge in a form of compositionality
 - Self-organization from a random start
- Parallel distributed processing
 - Tolerance of partially conflicting evidence
 - Tolerance of graceful degradation
 - Content-addressable memory
 - Low computational costs






DeePMD ([source](#))

A new HPC+DL+Physical paradigm

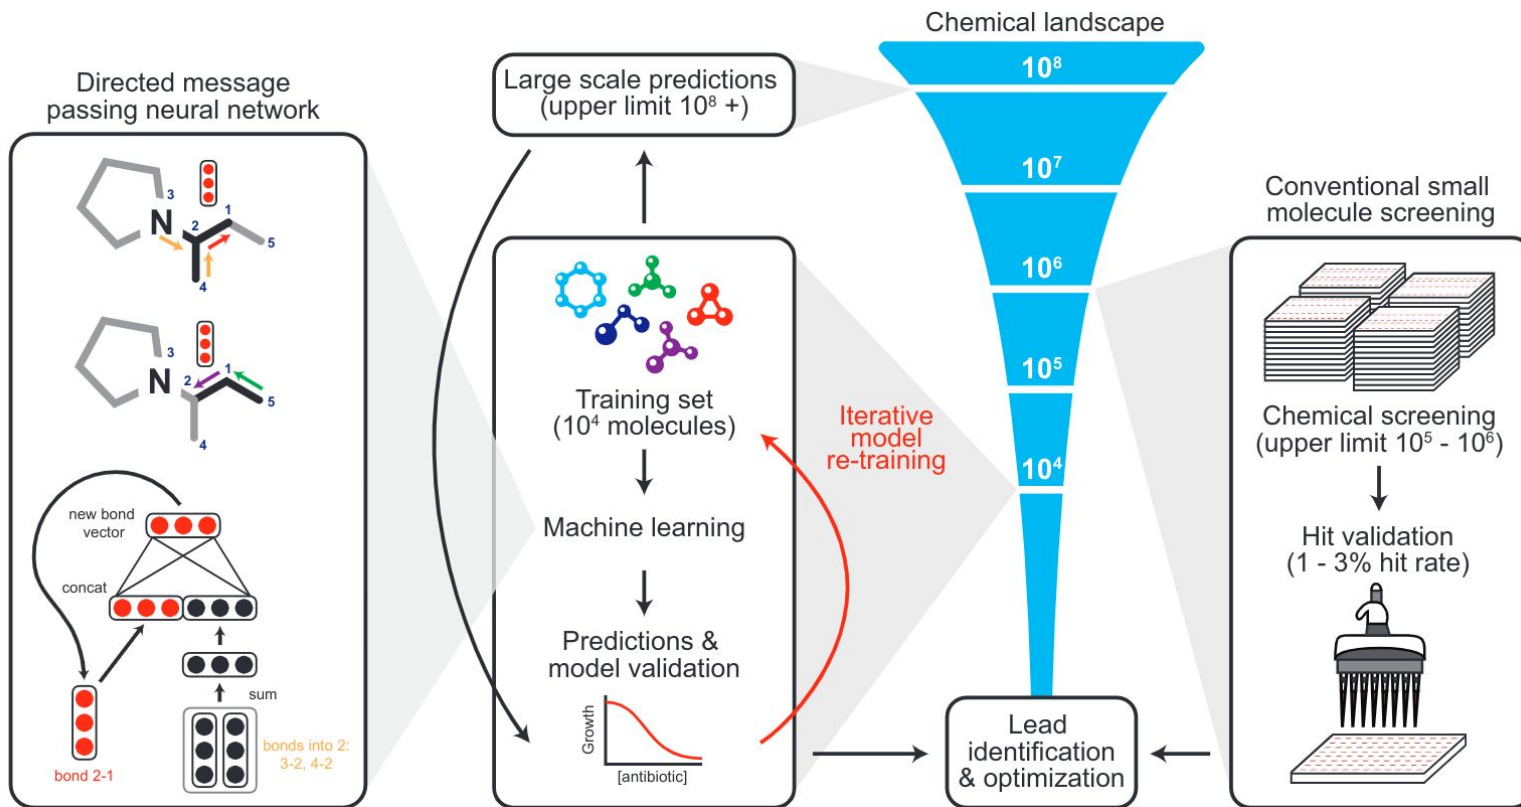


Pictures were from their SC'20 talk.



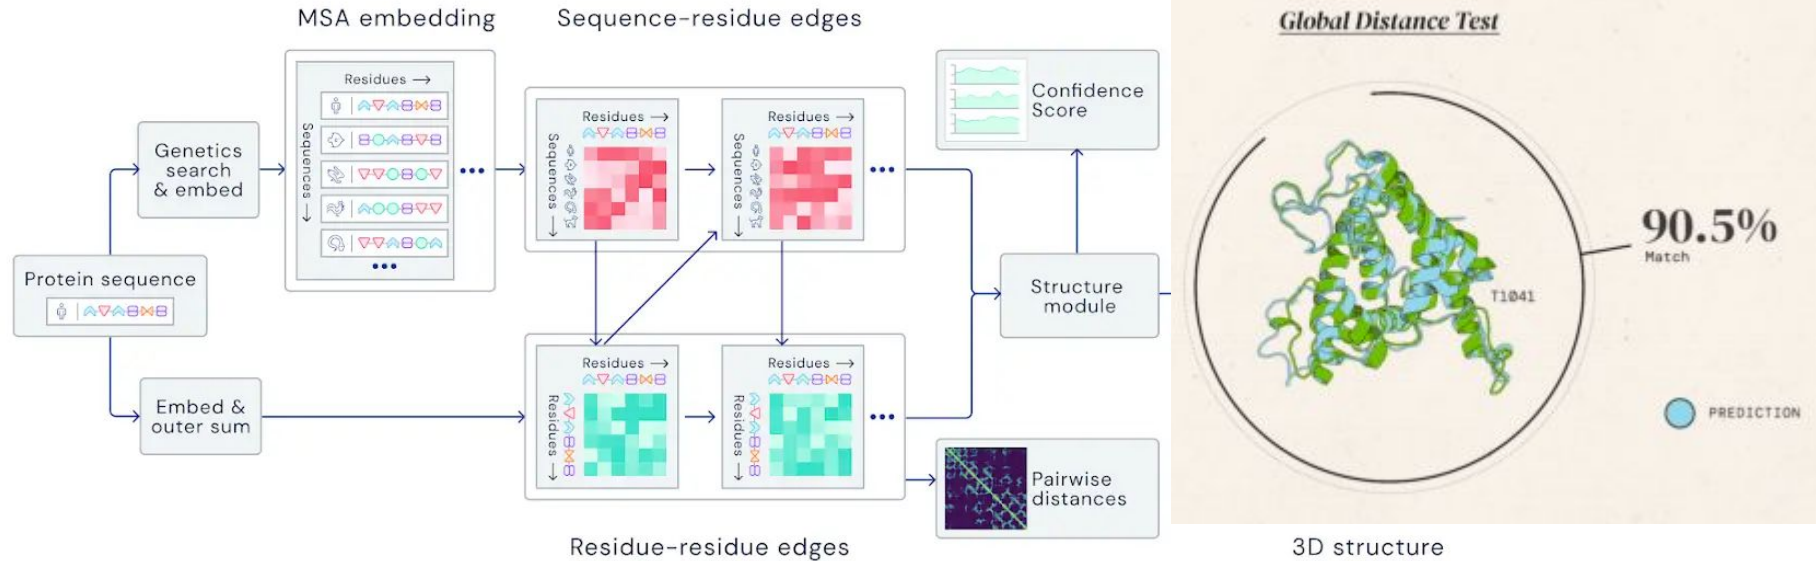


A deep learning approach to antibiotic discovery





AlphaFold: a significant breakthrough of AI on Science

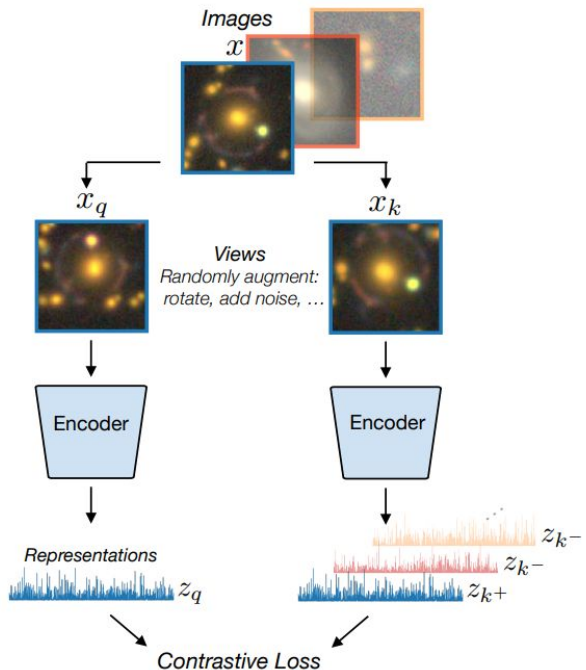


Astronomical discoveries through self-supervised learning




1. Self-supervised contrastive representation learning

Learn representations in an unsupervised manner



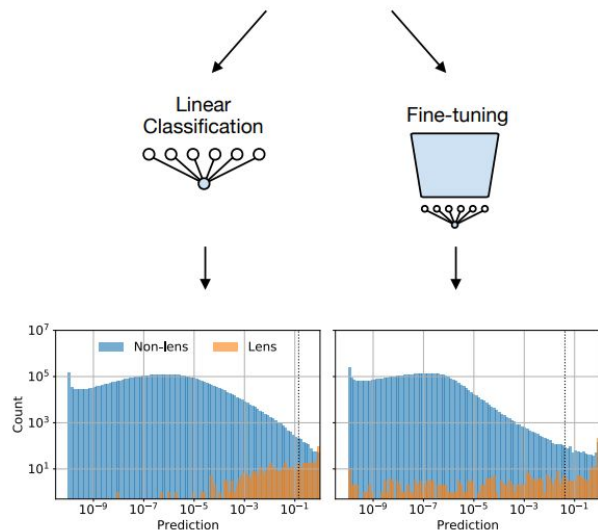
2. Similarity Search

Discover new lenses using a single labelled example



3. Automated classification

Introduce labels



<https://t.co/fFMnqHHJC0?amp=1>



John F Wu
@jwuphysics



I don't think it's an exaggeration to say that some of the biggest astronomical discoveries in the 2020s are only going to be made possible through self-supervised learning (SSL).




Driving forces and positive impacts

Data

- **ImageNet** database: Milestone of boosting Deep Learning
- **Kaggle**: a home of data scientists and analysts
- **Data explosion** from internet

Economy

- Venture Capital Investment soars **20x** in 8 years
- **Hundreds of million** dollars AI Merger & Acquisitions
- **100x** more people working on deep learning



Hardware

- AI compute amount increases **10 times** per year
- **GPU**, TPU, IPU, xPUs, Neuromorphic chips, Quantum AI chips...

Software

- From C++/Cuda to scripting languages (**Python**, R)
- From library packages to **frameworks**
- From toolsets to open-source pre-trained models

Neutral

General Intro



- What's ANN?
- DL, ML and AI

01

Good

DL as hopes



- Why DL is so powerful?
- Exciting achievements
- Mutual driving forces

02

Bad

DL as hypes



- Difficulties in application perspective
- Illusions in theoretical perspective

03

"Ugly"

DL like deep oceans



- Cracking the black box
- Thinking outside the box
- What about the future?

04



The more, the better?

Data, high-quality data!

"A child wearing sunglasses is labeled as a 'failure, loser, nonstarter, unsuccessful person' in the original ImageNet database ..."

-- Crawford K. et al. "Excavating AI"



- Quality of data determines the upper bound of how good the DL model can become.
- Tough tasks for data labeling and annotation
- Avoiding bias brings privacy concerns

Models grow way too big.

- GPT-1 (2018): 117 million parameters
- GPT-2 (2019): 1.5 billion parameters
- GPT-3 (2020): 175 billion parameters
- GPT-4 (2022): 100 trillion parameters



- Some fundamental limitations still there.
- Lower explainability in models.
- Algorithm efficiency is more important.
- Model's robustness is always a question.



The better, the better?

High expectations, low acceptance

Eye screening process in 11 clinics in Thailand

- Before DL evaluation: 2~10 weeks
- With DL evaluation: 10 mins

— Google Health, SIGCHI 2020



- Poor-equipped environment
- Unexpected corner cases
- Practical time/economical concerns

Where is DL in the pandemic?

"...none of them [papers] produced tools that would be good enough to use in a clinical setting. Something has gone seriously wrong when more than 300 papers are published that have no practical benefit."

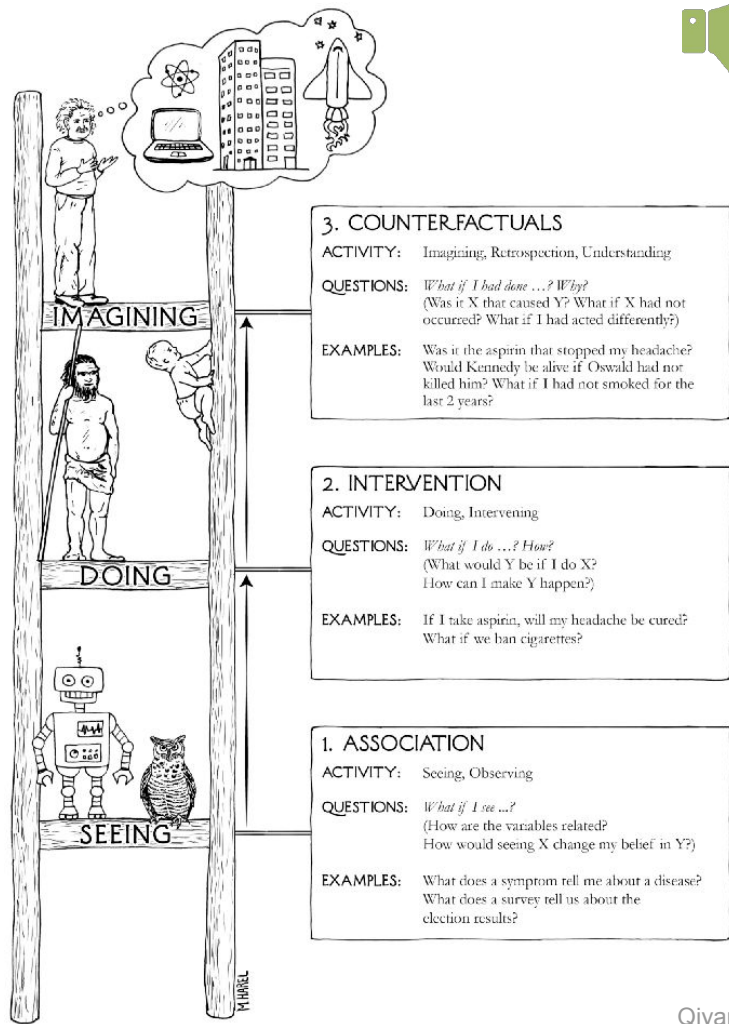
— "Nature Machine Intelligence" 2021



- Paper without details to reproduce results.
- Papers introduced significant biases.
- Publication bias towards positive results

Is it the real intelligence?

- Data mining \Rightarrow understanding?
 - Neural network = Curve fitting
 - Statistical = model-blind
- DL need something else to work together.
 - Causal reasoning: inductive \rightarrow deductive ?
 - Hybrid system with GOFAI? with Physics?
- A curse from Gödel's theorem
 - Maybe it is just impossible.
 - We don't even know what AI is/means.
- What *if* the Singularity finally happens
 - S-believers vs S-sceptics
 - Actions from AI-communities, policy makers



Neutral

General Intro



- What's ANN?
- DL, ML and AI

01

Good

DL as hopes



- Why DL is so powerful?
- Exciting achievements
- Mutual driving forces

02

Bad

DL as hypes



- Difficulties in application perspective
- Illusions in theoretical perspective

03

"Ugly"

DL like deep oceans



- Cracking the black box
- Thinking outside the box
- What about the future?

04




Understanding Neural Networks

Effective Theory

- NN \Rightarrow solving layer-to-layer iteration eqns + nonlinear learning dynamics
- Explaining nonlinearity to nth-order
- NN Predictions: Nearly Gaussian Dis. Depth-2-width ratio: deviation from ∞ -width Gaussian description
- Representation group flow: tuning the networks to criticality

Information Bottleneck

- NN: squeezing the information through a bottleneck
- Deep Learning proceeds 2 phases:
 - A short “fitting”
 - A much longer “compression”
- Network converges to information bottleneck theoretical bound



Discretized Dynamics

- Deep Network \Rightarrow Diff Eqns (DE)
- Network Architecture \Rightarrow Numerical DE
- Network Training \Rightarrow Optimal Control
- ResNet \Rightarrow One type of discretization
- Forward & Backprop \Rightarrow ODE solvers

Geometric Unification

- Symmetry and invariance
- Unify CNN, GNN, RNN, Transformer
- 5Gs: Grid, Groups, Graph/Manifold, Geodesics, Gauges
- A principled way to construct new NN frameworks


Analytical Philosophy

- Physical symbols: entire networks of cells, *not* locatable neurons
- Concepts: partially conflicting constraints, *not* strict logical definitions




Alternative Neural Networks

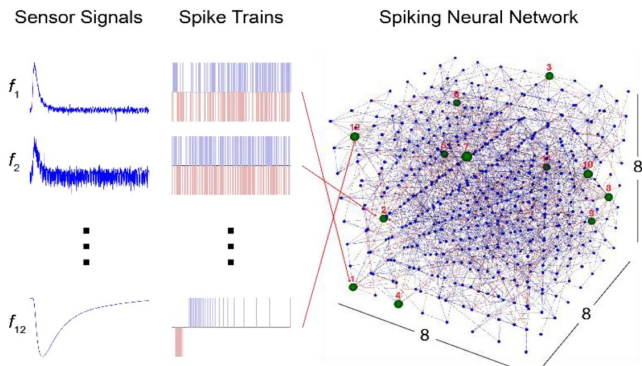
Capsule Network



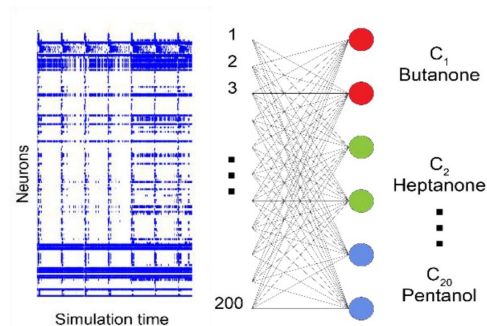
GasNets



Spiking Neural Network



Spiking Patterns Output Layer Classification





Hype Cycle for Artificial Intelligence, 2021



OARC Workshop Survey

<https://forms.gle/nbWgNP45qCwZhLRh9>