

Learning PyTorch

Qiyang Hu

UCLA Office of Advanced Research Computing

Nov 15, 2021



What is **PYTORCH**

- An open-source Python-based deep learning framework
 - Primarily developed by Facebook's AI Research lab (FAIR)
 - Replacement for Numpy with supporting GPUs, ROCm, TPUs
 - A full set of deep learning libraries
- History
 - Lua-based Torch (2002 - 2011): TH, THC, THNN, THCUNN
 - PyTorch 0.1 (2016): Python-based Torch
 - PyTorch 1.0 (2018): merging Caffe2
 - PyTorch 1.10 (Oct 21, 2021)
- PyTorch as a backend building block
 - Used in Tesla Autopilot, Uber's Pyro, Hugging Face's Transformers
 - Keras-like: PyTorch Lightning, PyTorch Ignite, tensorlayers, fast.ai
 - For specific domains: FlowTorch, NiftyTorch, Flair, Kornia, Skorch, ELF, Detectron2

Why **PYTORCH**

- **Simplicity**

- Feels like Numpy
- Consistent & great APIs

A graph is created on the fly

```
from torch.autograd import Variable  
  
x = Variable(torch.randn(1, 10))  
prev_h = Variable(torch.randn(1, 20))  
W_h = Variable(torch.randn(20, 20))  
W_x = Variable(torch.randn(20, 10))
```



- **Flexibility**

- Defining the model
- Modifying the model

- **Dynamic compute graphs**


- Immediate forward execution
- Tape-based autograd
- Destroyed immediately after backprop

- **Model serialization and quantization**

- JIT, TorchScript, FX
 - Seamlessly switch between Modes, Distributed training, Mobile deployment

Tensors as building blocks

1

$$\begin{bmatrix} 2 \\ 5 \\ 4 \end{bmatrix}$$
$$\begin{bmatrix} 2 & 1 & 3 \\ 5 & 7 & 9 \\ 4 & 8 & 6 \end{bmatrix}$$
$$\begin{bmatrix} \begin{bmatrix} 4 & 6 & 9 \\ 1 & 2 & 5 \\ 8 & 7 & 3 \end{bmatrix} \end{bmatrix}$$


Scalar
0-D

Vector
1-D

Matrix
2-D

Tensor
3-D

Tensor
 n -D

$X = 1$

$X[1] = 5$

$X[2, 1] = 8$

$X[0, 1, 2] = 5$


$X[\underbrace{2, 3, \dots, 1}_{N \text{ indices}}] = 6$

`torch.tensor([[[[1.0,1.0],[2.0,2.0]],[[3.0,3.0],[4.0,4.0]]],[[5.0,5.0],[6.0,6.0]],[[7.0,7.0],[8.0,8.0]]])`

Tensor, Storage and Views


$$M(i, j) = \text{offset} + \text{stride}[0] \cdot i + \text{stride}[1] \cdot j$$

- Data and Metadata




sizes	(D,H,W)
dtype	integer
device	cuda:0
layout	strided
strides	(H*W,W,1)

- Contiguous & Unboxed



- Size, offset, stride



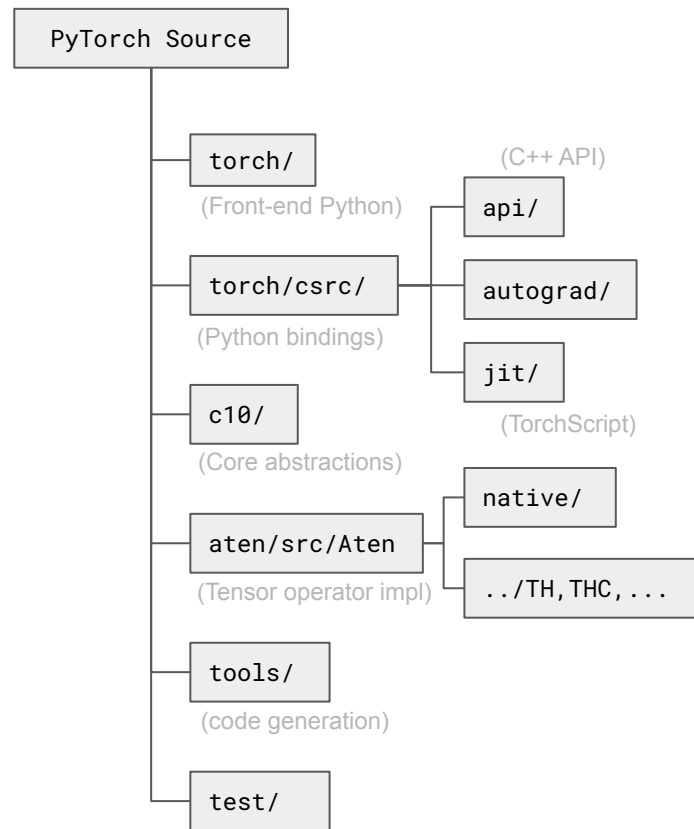
“Py” and “Non-Py” in PyTorch

- Tensor extensions

- Beyond strided tensors:
sparse, quantized, encrypted, MKLDNN, XLA tensors etc.
- Tensor wrapper: device ✖ layout ✖ dtype

- PyTorch = Python + C/C++ + CUDA

- Python extension objects in C/C++
- Code base components:
 - The core Torch libraries: TH, THC, THNN, THCUNN
 - Vendor libraries: CuDNN, NCCL
 - Python Extension libraries
 - Additional 3rd-party libraries: NumPy, MKL, LAPACK, DLPack



Colab Hands-on

bit.ly/learning_pytorch

Automatic differentiation

- Autograd package
 - Track all operations of tensors
 - Compute derivatives analytically via back-prop
 - Natively loaded in torch module
 - Can be used in other scientific domains
- Simple usage
 - Set tensor's `.requires_grad` as `TRUE`
 - Tensor's creation function recorded in `.grad_fn` attribute
 - Gradient accumulated into `.grad` attribute
 - Call `.backward()`
- Stop a tensor from tracking history
 - Wrap the code block in with `torch.no_grad()`
 - `.detach()`

Optimizers in PyTorch

- [torch.optim](#) package
 - Provides various optimization algorithms
 - Construct an optimizer object
 - `optimizer = optim.SGD(model.parameters(), lr=0.01)`
 - Need to move model to GPU before constructing optimizers
 - Must zero the gradient explicitly:
 - `optimizer.zero_grad()`
 - Take an optimization step:
 - `optimizer.step()` in GD method
 - `optimizer.step(closure)` in CG or LBFGS method
 - Optional: adjust the learning rate based on the number of epochs.
 - `optimizer.lr_scheduler`

Neural Networks in PyTorch

- [torch.nn](#) package
 - Contains all building blocks for neural network related work
 - `nn.functional` and `nn.Module`
- Define a network
 - For simple networks: concatenate modules through a `nn.Sequential` container
 - For complex networks: Subclassing `nn.Module`
- `nn.Module` package expects first index as batch size of samples
 - Need to reshape the input by `.unsqueeze()`
 - Use `Dataset` and `DataLoader`
- Loss functions in `torch.nn`:
 - `nn.MSELoss` (regression), `nn.BCELoss` (binary classification), `nn.CrossEntropyLoss` (multiclass classification)

PyTorch



Tensorflow

JAX

OARC Workshop Survey

<https://forms.gle/nbWgNP45qCwZhLRh9>