

Learning PyTorch

Qiyang Hu

UCLA Office of Advanced Research Computing

Nov 15, 2021



What is **PYTORCH**

- An open-source Python-based deep learning framework
 - Primarily developed by Facebook's AI Research lab (FAIR)
 - Replacement for Numpy with supporting GPUs, ROCm, TPUs
 - A full set of deep learning libraries
- History
 - Lua-based Torch (2002 - 2011)
 - PyTorch 0.1 (2016): THNN
 - PyTorch 1.0 (2018): merging Caffe2
 - PyTorch 1.10 (Oct 21, 2021)
- PyTorch as a backend building block
 - Used in Tesla Autopilot, Uber's Pyro, Hugging Face's Transformers
 - Keras-like: PyTorch Lightning, PyTorch Ignite, tensorlayers, fast.ai
 - For specific domains: NiftyTorch, Flair, Skorch, ELF, Detectron2

Why **PYTORCH**

- **Simplicity**

- Feels like Numpy
- Consistent & great APIs

A graph is created on the fly

```
from torch.autograd import Variable

x = Variable(torch.randn(1, 10))
prev_h = Variable(torch.randn(1, 20))
W_h = Variable(torch.randn(20, 20))
W_x = Variable(torch.randn(20, 10))
```



- **Flexibility**

- Defining the model
- Modifying the model

- **Dynamic compute graphs**

- Immediate forward execution
- Tape-based autograd
- Destroyed immediately after backprop

- **Model serialization and quantization**


- JIT, TorchScript, Tracing, FX
 - Seamlessly switch between Modes, Distributed training, Mobile deployment

Tensors as building blocks

1

$$\begin{bmatrix} 2 \\ 5 \\ 4 \end{bmatrix}$$

$$\begin{bmatrix} 2 & 1 & 3 \\ 5 & 7 & 9 \\ 4 & 8 & 6 \end{bmatrix}$$

$$\begin{bmatrix} \begin{bmatrix} 4 & 6 & 9 \\ 1 & 2 & 5 \\ 8 & 7 & 3 \end{bmatrix} \end{bmatrix}$$


Scalar
0-D

Vector
1-D

Matrix
2-D

Tensor
3-D

Tensor
 n -D

$$X = 1$$

$$X[1] = 5$$

$$X[2, 1] = 8$$


$$X[0, 1, 2] = 5$$

$$X[\underbrace{2, 3, \dots, 1}_{N \text{ indices}}] = 6$$



Tensor, Storage and Views

$$M(i, j) = \text{offset} + \text{stride}[0] \cdot i + \text{stride}[1] \cdot j$$

- Data and Metadata
- Contiguous & Unboxed
- Offset, size, stride

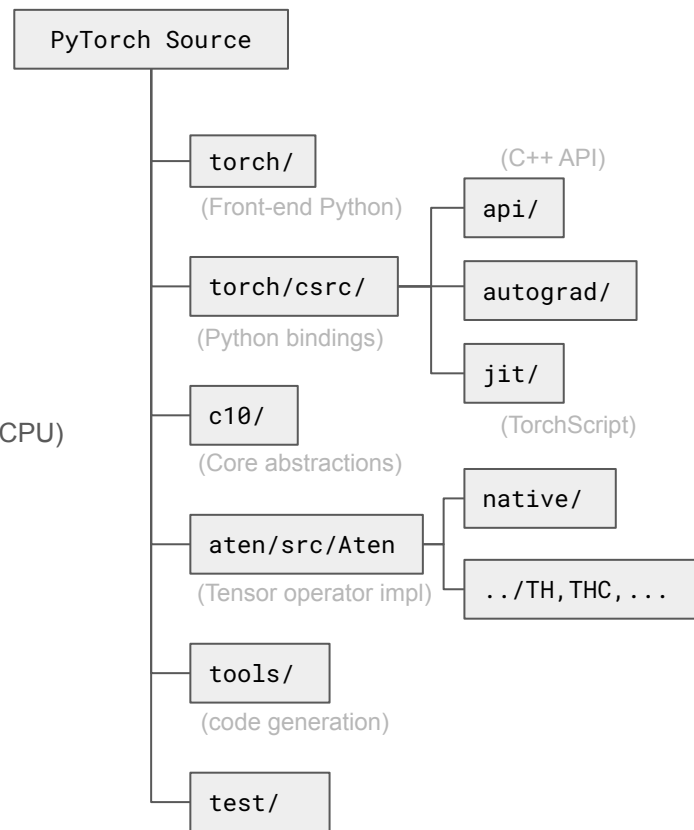


sizes	(D,H,W)
dtype	integer
device	cuda:0
layout	strided
strides	(H*W,W,1)



“Py” and “Non-Py” in PyTorch

- Tensor extensions
 - Beyond strided tensors:
sparse, quantized, encrypted, MKLDNN, TPU tensors etc.
 - Tensor wrapper: device ✕ layout ✕ dtype
- Works with Numpy arrays
 - Easy conversion
 - Zero copy: share their underlying memory locations (if on CPU)
- PyTorch = Python + C/C++ + CUDA
 - Python extension objects in C/C++
 - Code base components:
 - The core Torch libraries: TH, THC, THNN, THCUNN
 - Vendor libraries: CuDNN, NCCL
 - Python Extension libraries
 - Additional 3rd-party libraries: NumPy, MKL, LAPACK, DLPack



Colab Hands-on

bit.ly/learning_pytorch

Automatic differentiation

- Autograd package
 - Track all operations of tensors
 - Compute derivatives analytically via back-prop
 - Natively loaded in torch module
 - Can be used in other scientific domains
- Simple usage
 - Set tensor's `.requires_grad` as `TRUE`
 - Call `.backward()`
 - Gradient accumulated into `.grad` attribute
 - Tensor's creation function recorded in `.grad_fn` attribute
- Stop a tensor from tracking history
 - `.detach()`
 - Wrap the code block in with `torch.no_grad()`

Optimizers in PyTorch

- [torch.optim](#) package
 - Provides various optimization algorithms
 - Need to move model to GPU before constructing optimizers
 - Must zero the gradient explicitly:
 - `optimizer.zero_grad()`
 - Take an optimization step:
 - `optimizer.step()` in GD method
 - `optimizer.step(closure)` in CG or LBFGS method
 - Optional: adjust the learning rate based on the number of epochs.
 - `optimizer.lr_scheduler`

Neural Networks in PyTorch

- [torch.nn](#) package
 - Contains all building blocks for NN architectures
 - All blocks subclassed from `nn.Module` (e.g. `nn.Linear`)
- Define a network
 - For simple networks: concatenate modules through a `nn.Sequential` container
 - For complex networks: Subclassing `nn.Module`
- `nn.Module` package expects first index as first batch size of samples
 - Need to reshape the input by `.unsqueeze()`
 - Use `Dataset` and `DataLoader`
- Loss functions in `torch.nn`:
 - `MSELoss` (regression), `BCELoss` (binary classification), `CrossEntropyLoss` (multiclass classification)

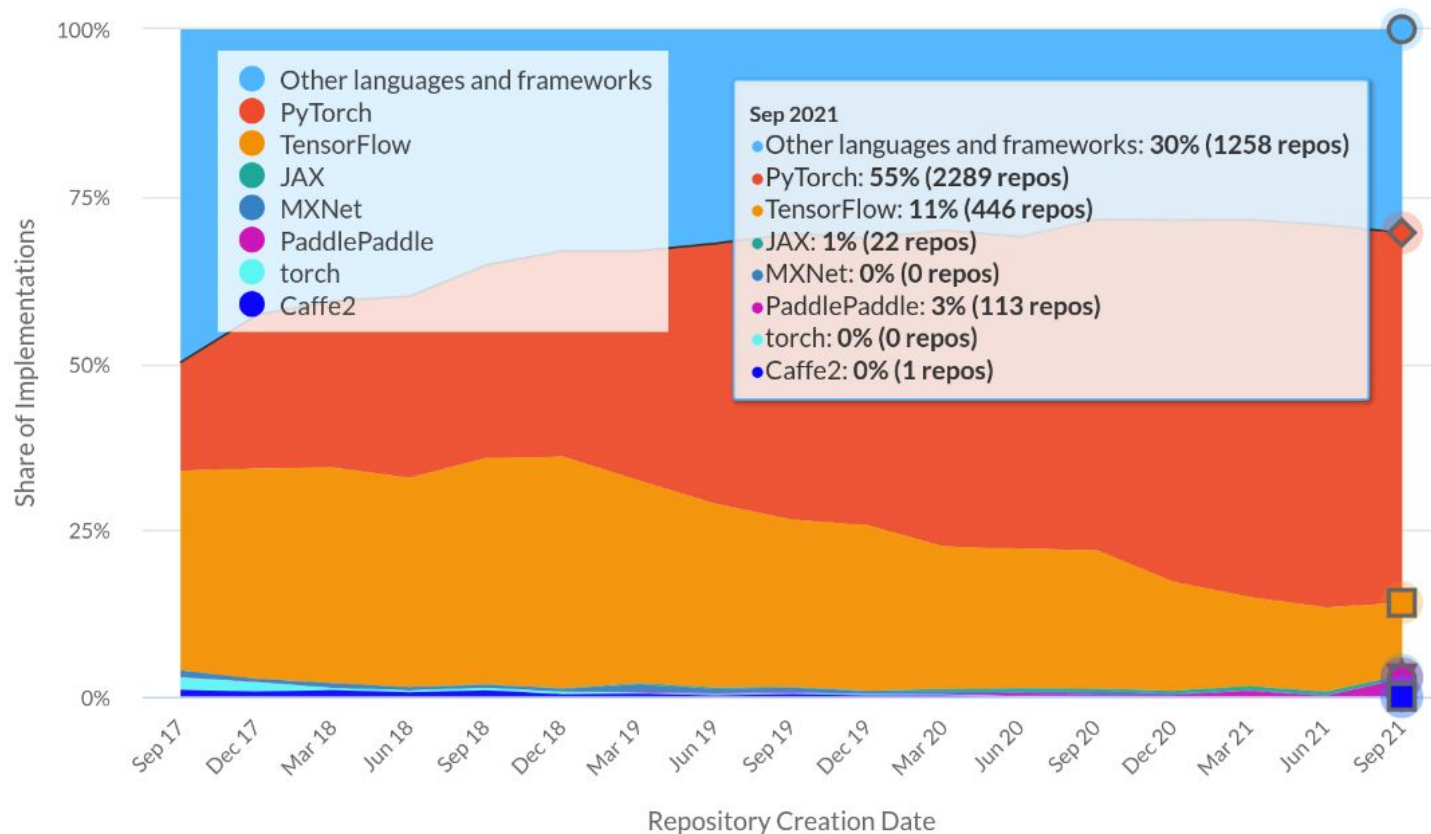
PyTorch



Tensorflow

JAX

Frameworks used in paper implementations on GitHub



OARC Workshop Survey

<https://forms.gle/nbWgNP45qCwZhLRh9>