

# **Learning Convolutional Neural Networks (2)**

Qiyang Hu

UCLA Office of Advanced Research Computing

Feb 18, 2022

# In this talk

[bit.ly/LDL\\_repo](https://bit.ly/LDL_repo)

## Improving our model



- Data Augmentation
- Dropout & batch norm
- Demo

01

## Transfer learning



- Transferring knowledge
- MobileNet
- Demo

02

## Latest Developments



- ViT & Swin-T
- ConvViT & ConvNeXt

03

# In this talk

## Improving our model



- Data Augmentation
- Dropout & batch norm
- Demo

01

## Transfer learning



- Transferring knowledge
- MobileNet
- Demo

02

## Latest Developments





- ViT & Swin-T
- ConvViT & ConvNeXt

03

# Quick Recap

- Dogs-vs-Cats challenges
  - 25,000 training images
  - 15,000 testing images
- Construct our own CNNs
  - 4 Conv layer blocks
  - Flatten layer
  - Dense layer
- Overfitting
  - Memorizing training set too much
  - Missing the essence knowledge
- How to improve?
  - Need more training data
  - Need regularization

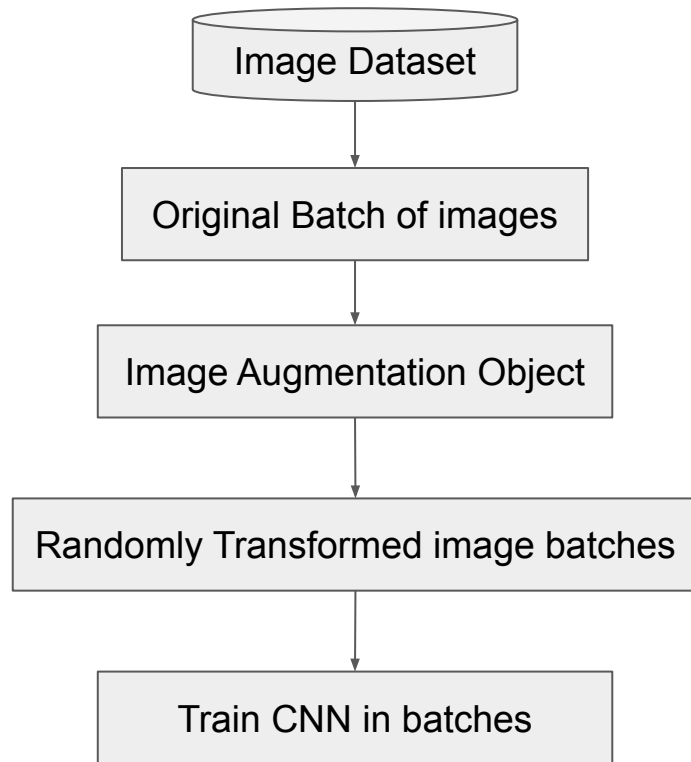


# Data and models: the bigger the better, *really*?

	VGGNet	DeepVideo	GNMT	GPT-3
Used For	Identifying Image Category	Identifying Video Category	Translation	Text Generation
Input	Image	Video	English Text	Text
Output	1000 Categories	47 Categories	French Text	Text
Parameters	140M	~100M	380M	175B
Data Size	1.2M Images with assigned Category	1.1M Videos with assigned Category	6M Sentence Pairs, 340M Words	2TB of Internet Text ~ 499 Billion Tokens

# From big size to smart size

- Using data augmentation
  - To get more data with “no more”
  - Various transformations to the available dataset
  - Prevent the irrelevant data
- Types of data augmentation
  - Offline augmentation
    - Performing all the transformations beforehand
    - Good for smaller dataset
  - In-place augmentation
    - Performing transformations in mini-batches
    - Preferred for larger dataset
- Data augmentation in PyTorch
  - torchvision.transforms




# Augmentation Techniques

- Flip
- Affine Transformation
  - Rotation
  - Zoom & Crop
  - Translation
- Gaussian Noise
- ZCA whitening
- Histogram Equalization
- Feature-wise standardization
- Neural Style Transfer

Input Image


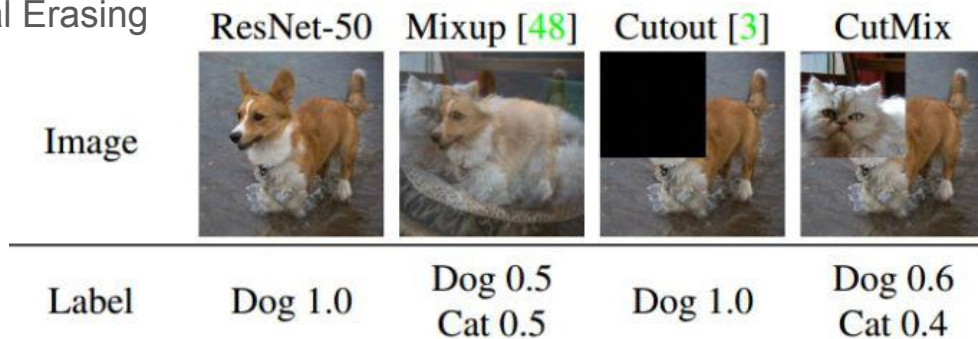


Augmented Images



# More Data Augmentation Techniques in CV tasks


- Random erasing: `torchvision.transforms.RandomErasing(p=0.5, scale=(0.02, 0.33), ...)`
  - Cutout (masking out random sections): no label change
  - Hide-and-seek, GridMask
  - Object Region Mining with Adversarial Erasing
- Mixup: soft overlapping
- Cutmix/Mosaic: hard masking
  - Cutmix: 2 images mixed
  - Mosaic: 4 images mixed
- 3-D augmentation






# Drop-out technique

- Gradient vanishing during DNN training:
  - Imbalanced weights in network:
    - Larger weights => well trained
    - Smaller weights => not trained that much!
- Dropout: randomly turns off some neurons
  - Forcing networks to train weak neurons
  - Dropout rate: default 50%
    - Roughly double the iterations to converge
  - Training time in epoch is less
  - Srivastava 2014 [paper](#)
  - Variations: spatial dropout, etc.
- PyTorch: [torch.nn.Dropout2d\(\)](#)
  - Implemented by “Inverted dropout” technique
  - Apply to the corresponding layer(s)




(a) Standard Neural Net



(b) After applying dropout.

# Other training techniques in deep learning

- Regularizer
  - L1(Lasso), L2(Ridge), L1\_L2(ElasticNet) in each layer
    - `weight_decay` flag for L2 in pytorch optimizers
- Early Stopping
  - Stop training when validation loss reach minimum
- Batch Normalization
  - Normalize the data (input features) across batches in each mini-batch
  - Add batch normalization before activation function
    - `torch.nn.BatchNorm2D(num_features=n_chans1)`
- Skip connections
  - Simple trick to add the input (conv1) to the output of a block of layers (conv3)
  - Residual networks ([K. He, 2015](#))
    - Opened the door to hundreds-layer-depth networks (Highway Net, U-Net, Dense-Nets, ...)



# Colab Hands-on

[bit.ly/LDL\\_cnn2](https://bit.ly/LDL_cnn2)

# Tips of CNNs

- Design

- Kernels

- Shape: square for visual tasks, rectangles for NLPs
    - Size: odd x odd, (3x3, 5x5, 7x7)

- Number of Conv Layers:

- Start with a few
    - < 100

- Hyperparameters:

- Batch size: neither too big, nor too small


- start from 16, then 8 or 32

- Learning rate: start from 0.01

- minibatch size \* k  $\Rightarrow$  learning rate \* k

- Use batch normalization

- Use random search



[Source](#)

# In this talk

## Improving our model



- Data Augmentation
- Dropout & batch norm
- Demo

01

## Transfer learning



- Transferring knowledge
- MobileNet
- Demo

02

## Latest Developments




- ViT & Swin-T
- ConvViT & ConvNeXt

03

# Transfer Learning

- Reusing the developed neural networks
  - Greatly speed up our training
  - Make it mobile
- Reused model  $\Rightarrow$  feature extractor
  - Pre-trained on a popular generic dataset
  - Transfer the knowledge
    - Match the input
    - Add new layers for specific data and tasks
- Two trends of pre-trained models
  - Gigantic backbones: ResNet, BERT
  - Light-fast-efficient backbones:
    - SqueezeNet, MobileNet v1, v2
    - MobileNet v3, TinyYOLO, MnasNet, ShuffleNet, CondenseNet, ESPNet, DiCENet, MixNet, EfficientNet


Should freeze  
or not freeze?



# MobileNet

- Very efficient CNNs ([v2 paper](#) & [v3 paper](#))
  - Depthwise separable convolutions
  - Inverted residual with linear bottleneck
  - Squeeze-and-excitation (SE) modules
- Loading the model:
  - MobileNet v2:
    - [torchvision.models](#)
    - [PyTorch Hub](#)
  - MobileNet v3:
    - large, small and quantized
    - `torchvision.models.mobilenet_v3_...`
- Modify the classifier layer

```
model.classifier[3] = torch.nn.Linear(1280, 2)
```



(MACs = Multiply-Accumulates)

Figure from [v2 paper](#) and [v3 paper](#)

# In this talk

## Improving our model



- Data Augmentation
- Dropout & batch norm
- Demo

01

## Transfer learning



- Transferring knowledge
- MobileNet
- Demo

02

## Latest Developments



- ViT & Swin-T
- ConvViT & ConvNeXt


03




# A Twist of Modern ConvNets

- Renaissance since 2012
  - AlexNet, VGGNet, ResNet, MobileNet, EfficientNet, ...
  - Rapid development due to
    - “Sliding window”
    - Translation equivariance
- Challenges from NLP (2020)
  - Vision Transformers (ViT), Swin Transformers (Swin-T)
  - Replacing ConvNet by:
    - Split img to a seq of patches
    - Permutation invariant via self-att
- Revival since 2021
  - ConViT (2021), [ConvNeXt \(2022\)](#)


Swin Transformer




ResNet



ConvNeXt



ImageNet-1K Acc.



# Survey

bit.ly/**survey\_cnn2**