

Learning Generative Adversarial Networks

Qiyang Hu

UCLA Office of Advanced Research Computing

Feb 25, 2022

In this talk

bit.ly/LDL_repo

Introduction



- What is GANs?
- Training of GANs

01

GANs Hands-On



- Today's project
- DCGANs
- GANs in PyTorch

02

GANs World



- Variant GANs
 - SAGAN, CGAN, CycleGAN
- Challenges in GANs

03

In this talk

Introduction



- What is GANs?
- Training of GANs

01

GANs Hands-On



- Today's project
- DCGANs
- GANs in PyTorch

02

GANs World



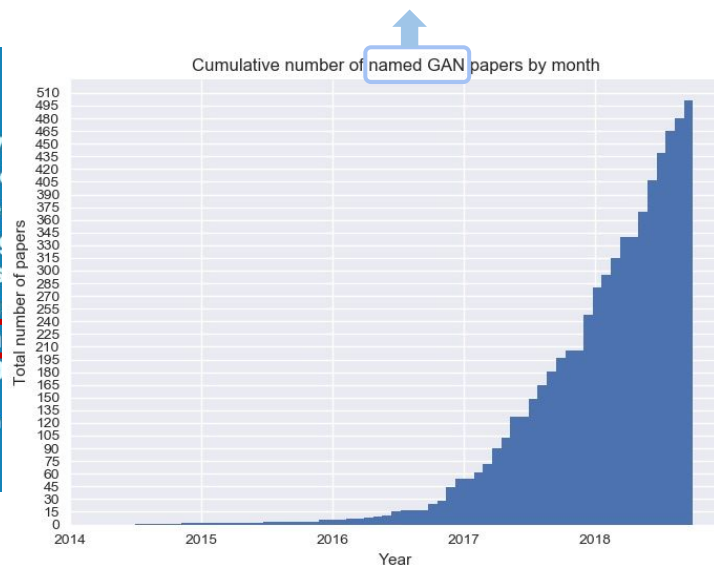
- Variant GANs
 - SAGAN, CGAN, CycleGAN
- Challenges in GANs

03



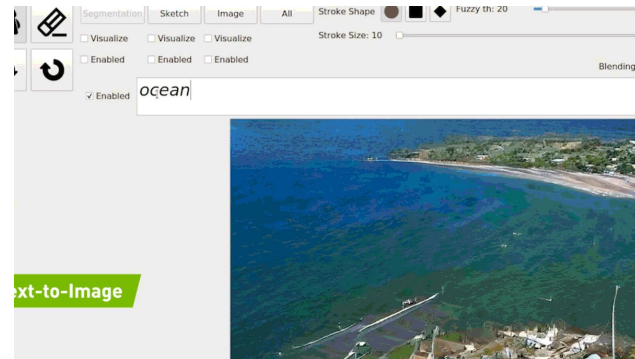
Generative Adversarial Networks



GAN, CGAN, DCGAN, RGAN, SGAN...



The only limit of GAN is our imagination



GAN: Generative

Training Data ([source](#))



• • • • • • • • • •

$$p_{\text{data}}(\mathbf{x})$$



Generated Samples ([source](#))




• • • • • • • • • •



$$p_{\text{model}}(\mathbf{x})$$

- Discriminative model


- discriminate labels of data instances
- try to draw boundaries in the data space



- capture the conditional prob. $p(Y | X)$
- measure the misfit of points
- learn the difference, ignore correlations

- Generative model

- generate new data instances
- try to model how data is placed



- capture the joint prob. $p(X, Y)$
- measure the misfit of prob distributions
- learn distributions to capture correlations

VS

Taxonomy of Generative Models

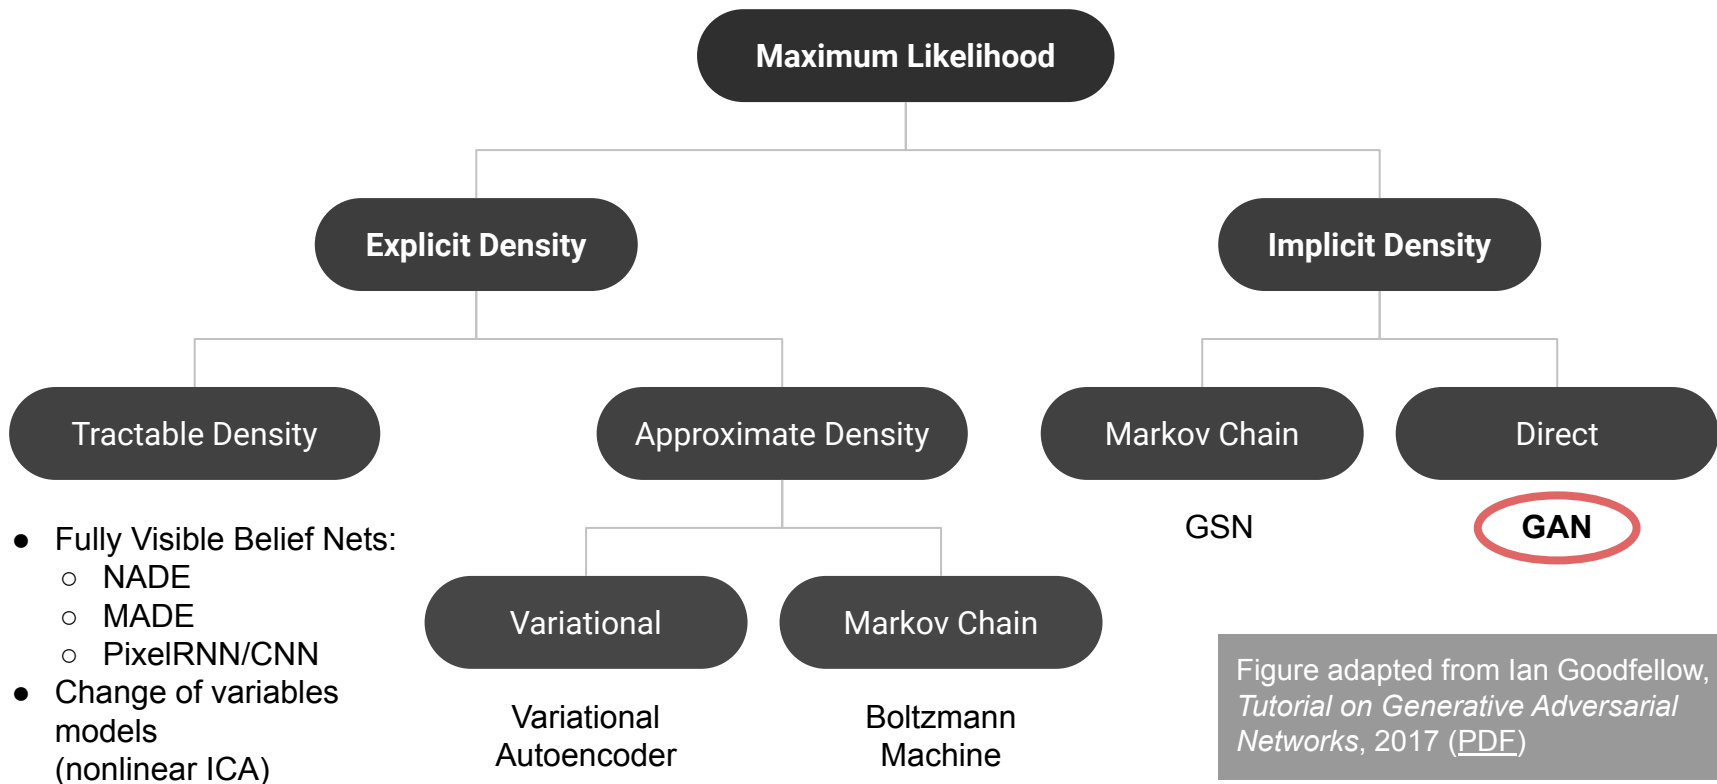



Figure adapted from Ian Goodfellow,
*Tutorial on Generative Adversarial
Networks*, 2017 ([PDF](#))

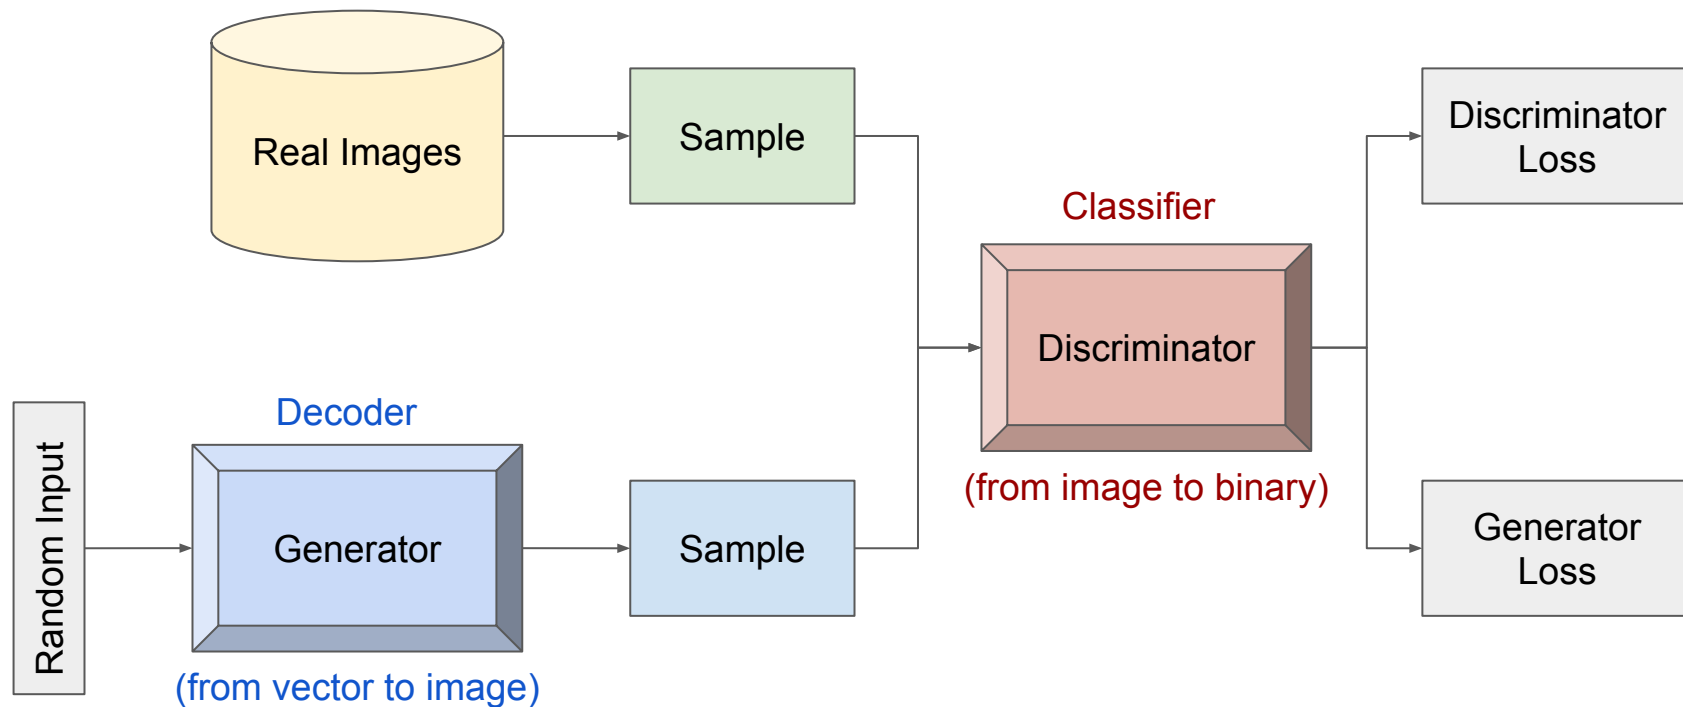
GAN: Adversarial

- **Generator**: generate plausible data
- **Discriminator**: distinguish the generator's fake data from real data

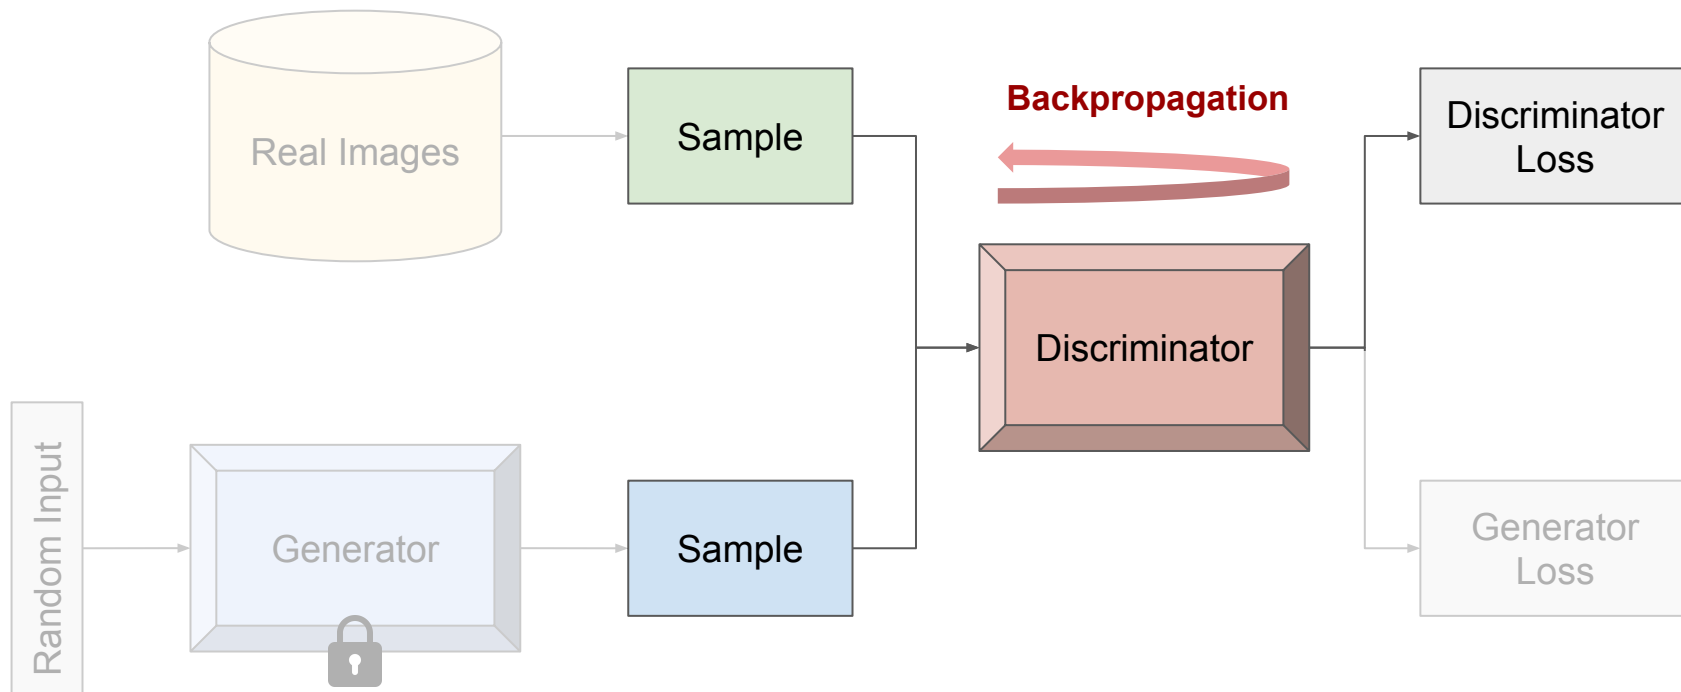


[Image Source](#)

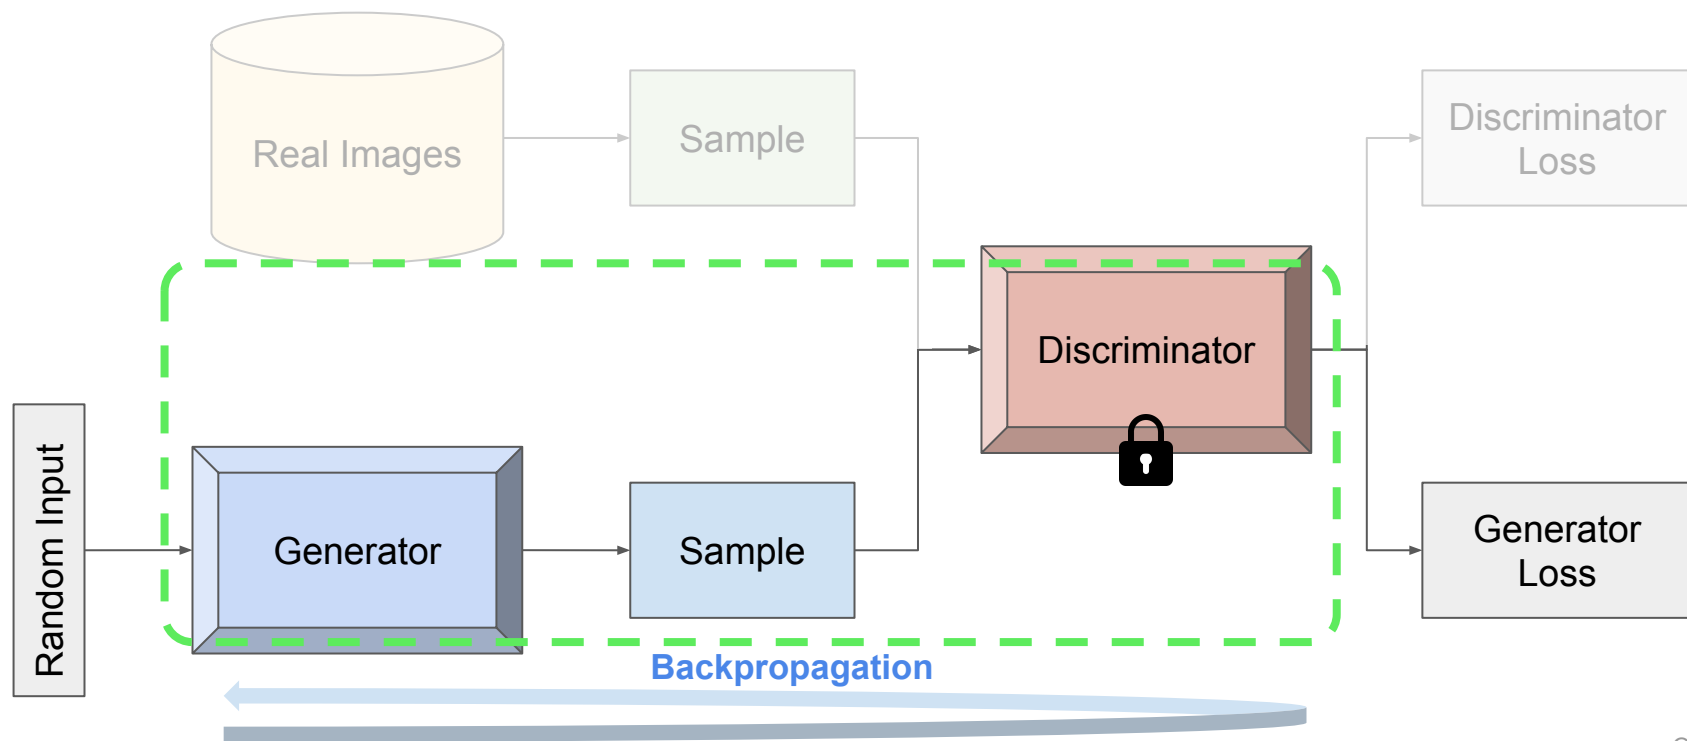
GAN: Network



Training of GAN (1): update discriminator



Training of GAN (2): update generator



Training of GAN (3): iterate the 2 steps to converge

- Alternate the training periods
 - The discriminator trains for one or more epochs with locking generator
 - The generator trains for one or more epochs with locking discriminator
 - Repeat the above steps
- When to stop
 - While generator improves, discriminator performance gets worse
 - Generator becomes perfect \Rightarrow discriminator gets 50% accuracy
 - Feedback from discriminator is less meaningful over time
 - At some point discriminator starts giving completely random feedback
 - Generator starts to train on junk feedback, and its own quality may collapse
 - Convergence of GANs is unstable, very hard to identify

The design logic behind the GAN structure

- Why do we need discriminator ?
 - There are generative models that can learn without discriminator
 - e.g. Variational Autoencoder (VAE)
 - Generator constructs the images in a bottom-up way
 - Very hard to capture the higher-level correlations
 - The discriminator can guide the generator with correlation info in a criticizing way
- Why do we need generator?
 - There are generative models that can learn without generator
 - e.g. Energy based model
 - Discriminator constructs the images in a top-down way
 - Very hard to learn from constructing negative sampling
 - The generated instances become negative training examples for the discriminator.

Standard Loss function for GAN

- Minimax Loss

- Proposed in the original [Goodfellow's paper](#)

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} (\log D(x)) + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

Real image sample

Noise sample

Fake image sample

Probability of the real image is real

Probability of the fake image is real

- Derives from a single measure of distance ([BCE](#)) between the real and generated distributions

- In practice

- Discriminator loss: maximize $\frac{1}{n} \sum_{i=0}^n \log(D(x_i)) + \frac{1}{n} \sum_{i=0}^n \log[1 - D(G(z_i))]$
- Generator (not-saturating) loss: maximize $\frac{1}{n} \sum_{i=0}^n \log[D(G(z_i))]$

GANs are *very* difficult to train.

- Discriminator shouldn't be too good.
 - Good discriminator \Rightarrow always 100% accuracy
 - Generator has no positive case to follow for learning.
 - Mathematically, falling into the vanishing-gradient zone
 - Generator needs some success, esp. in early stages
- Discriminator shouldn't be too bad.
 - Bad discriminator \Rightarrow random guess
 - Generator cannot get helpful feedback, esp. in late stages

Training Tips For GANs

- <https://github.com/soumith/ganhacks>
- Need experiences as always



Input Normalizing

- normalize the images to $(-1, 1)$
- Tanh as the last layer of the generator output

Implemented in Demo: Yes



Tune the learning rates

- Make D not improve too fast
- Make D not improve too slow

Implemented in Demo: Yes



A modified loss function

- Generator loss function to be $\max \log(D)$
- Flip labels when training generator:
 $real = fake, fake = real$

Implemented in Demo: Yes



BatchNorm

Construct different mini-batches for real and generated samples

Implemented in Demo: Yes



Add noise to inputs

- Perturb the both real and fake images when training D
- Decay the noise over time.

Implemented in Demo: No

...



Avoid Sparse Gradients: ReLU, MaxPool

- LeakyReLU is good for G and D
- Use stride, not pooling

Implemented in Demo: Yes



Use Soft and Noisy Labels

- Real $\sim \text{Uniform}(0.7, 1.2)$
- Fake $\sim \text{Uniform}(0.0, 0.3)$

As homework



Use DCGAN or Hybrid

- Use DCGAN if possible
- If not, use hybrid of KL + GAN or VAE + GAN

Implemented in Demo: Yes

ADAM

Use the ADAM Optimizer

optim.Adam rules.

Implemented in Demo: Yes



Use Dropouts in G

- Provide noise in the form of dropout (50%)
- Apply at both training and test time

Implemented in Demo: No

In this talk

Introduction



- What is GANs?
- Training of GANs

01

GANs Hands-On



- Today's project
- DCGANs
- GANs in PyTorch

02

GANs World



- Variant GANs
 - SAGAN, CGAN, CycleGAN
- Challenges in GANs

03

Today's Demo — Generative Dog Images from Kaggle

- Experiment with creating puppy pics
 - A Kernels-only competition (total 10K prize, expired years ago)
 - [Evaluation](#)
 - Using a pre-trained model (Inception)
 - Calculating MiFID scores
- Using [Stanford Dogs Dataset](#)
 - 20,580 images with annotation info (120 breeds, bounding box)
 - Some dog pictures are very *tricky*
 - Only part of the dogs body
 - Having multiple dogs
 - Having multiple persons
 - Dogs may occupy $< \frac{1}{5}$ of the picture
 - With various texts (from memes, magazine, etc)
 - Even wild predators included




Neural Networks for Generator and Discriminator

- Deep Convolutional Generative Adversarial Networks ([DCGANs](#))

- Generator

- Input: a std-norm latent vector
- Strided 2D Convolutional-transpose layers
- Batch norm layers
- ReLU activations
- Convtrans+*Tanh* before output
- Output: a 3x64x64 RGB image



- Discriminator

- Input: 3x64x64 input image
- Strided convolution layers, batch norm layers, LeakyReLU activations
- Conv+Sigmoid before output
- Output: a scalar probability

Generator Implementation in PyTorch

```
class Generator(nn.Module):
    def __init__(self, nz=128, channels=3):
        super(Generator, self).__init__()

        self.nz = nz
        self.channels = channels

    def convlayer(n_input, n_output, k_size=4, stride=2, padding=0):
        block = [
            nn.ConvTranspose2d(n_input, n_output, kernel_size=k_size, stride=stride, padding=padding, bias=False),
            nn.BatchNorm2d(n_output),
            nn.ReLU(inplace=True),
        ]
        return block

    self.model = nn.Sequential(
        *convlayer(self.nz, 1024, 4, 1, 0), # Fully connected layer via convolution.
        *convlayer(1024, 512, 4, 2, 1),
        *convlayer(512, 256, 4, 2, 1),
        *convlayer(256, 128, 4, 2, 1),
        *convlayer(128, 64, 4, 2, 1),
        nn.ConvTranspose2d(64, self.channels, 3, 1, 1),
        nn.Tanh()
    )

    def forward(self, z):
        z = z.view(-1, self.nz, 1, 1)
        img = self.model(z)
        return img
```

Discriminator Implementation in PyTorch

```
class Discriminator(nn.Module):
    def __init__(self, channels=3):
        super(Discriminator, self).__init__()

        self.channels = channels

    def convlayer(n_input, n_output, k_size=4, stride=2, padding=0, bn=False):
        block = [nn.Conv2d(n_input, n_output, kernel_size=k_size, stride=stride, padding=padding, bias=False)]
        if bn:
            block.append(nn.BatchNorm2d(n_output))
        block.append(nn.LeakyReLU(0.2, inplace=True))
        return block

    self.model = nn.Sequential(
        *convlayer(self.channels, 32, 4, 2, 1),
        *convlayer(32, 64, 4, 2, 1),
        *convlayer(64, 128, 4, 2, 1, bn=True),
        *convlayer(128, 256, 4, 2, 1, bn=True),
        nn.Conv2d(256, 1, 4, 1, 0, bias=False), # FC with Conv.
    )

    def forward(self, imgs):
        logits = self.model(imgs)
        out = torch.sigmoid(logits)

        return out.view(-1, 1)
```

Training loop

```
#####  
# (1) Update D network: maximize  $\log(D(x)) + \log(1 - D(G(z)))$   
#####  
# train with real  
netD.zero_grad()  
real_images = real_images.to(device)  
batch_size = real_images.size(0)  
labels = torch.full((batch_size, 1), real_label, device=device)  
  
output = netD(real_images)  
errD_real = criterion(output, labels)  
errD_real.backward()  
D_x = output.mean().item()  
  
# train with fake  
noise = torch.randn(batch_size, nz, 1, 1, device=device)  
fake = netG(noise)  
labels.fill_(fake_label)  
output = netD(fake.detach())  
errD_fake = criterion(output, labels)  
errD_fake.backward()  
D_G_z1 = output.mean().item()  
errD = errD_real + errD_fake  
optimizerD.step()
```

real_label $\neq 1$ to make
the discriminator not
learn too quickly

```
real_label = 0.9  
fake_label = 0
```

```
#####  
# (2) Update G network: maximize  $\log(D(G(z)))$   
#####  
netG.zero_grad()  
labels.fill_(real_label) # fake labels are real for generator cost  
output = netD(fake)  
errG = criterion(output, labels)  
errG.backward()  
D_G_z2 = output.mean().item()  
optimizerG.step()
```

Colab Hands-on

bit.ly/LDL_gan

In this talk

Introduction



- What is GANs?
- Training of GANs

01

GANs Hands-On



- Today's project
- DCGANs
- GANs in PyTorch

02

GANs World



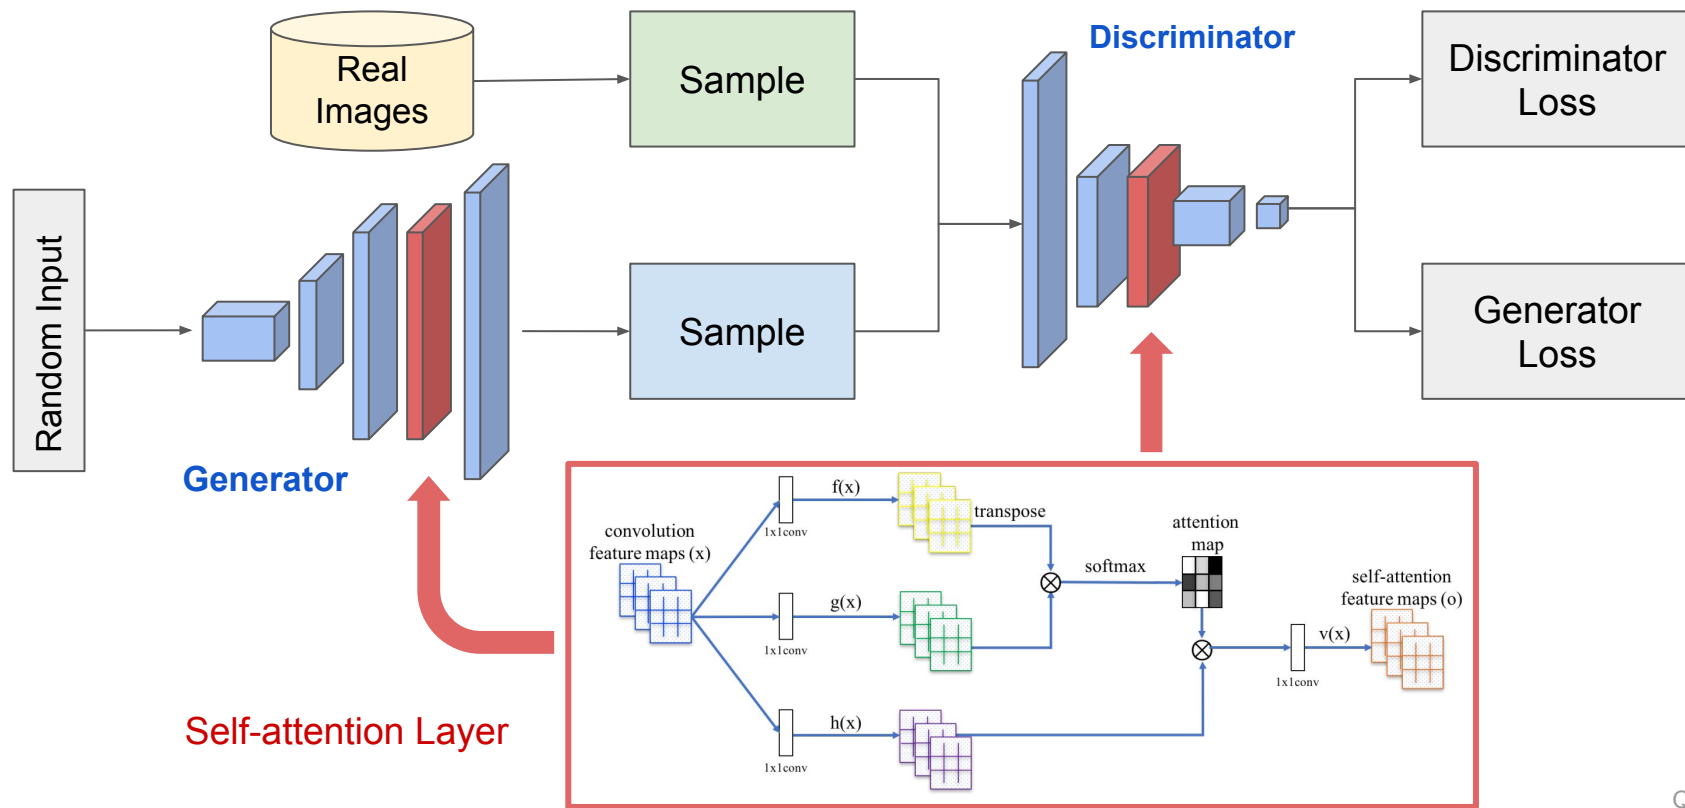
- Variant GANs
 - SAGAN, CGAN, CycleGAN
- Challenges in GANs

03


A lot of different GANs!

- Various design of network structures
 - SAGANs
 - Conditional GANs
 - CycleGANs
 - InfoGANs
 - EB-GANs
 - VAE-GANs
 - BiGANs
 - Triple-GANs
 - ...
- Various metrics for objective functions
 - WGANs
 - LSGANs
 - RGANs
 - Cramer GANs
 - Fisher GANs
 - MMD GANs
 - McGANs
 - HingeGANs
 - ...
- Combining the two
 - BEGAN
 - MAGANs
 - ...

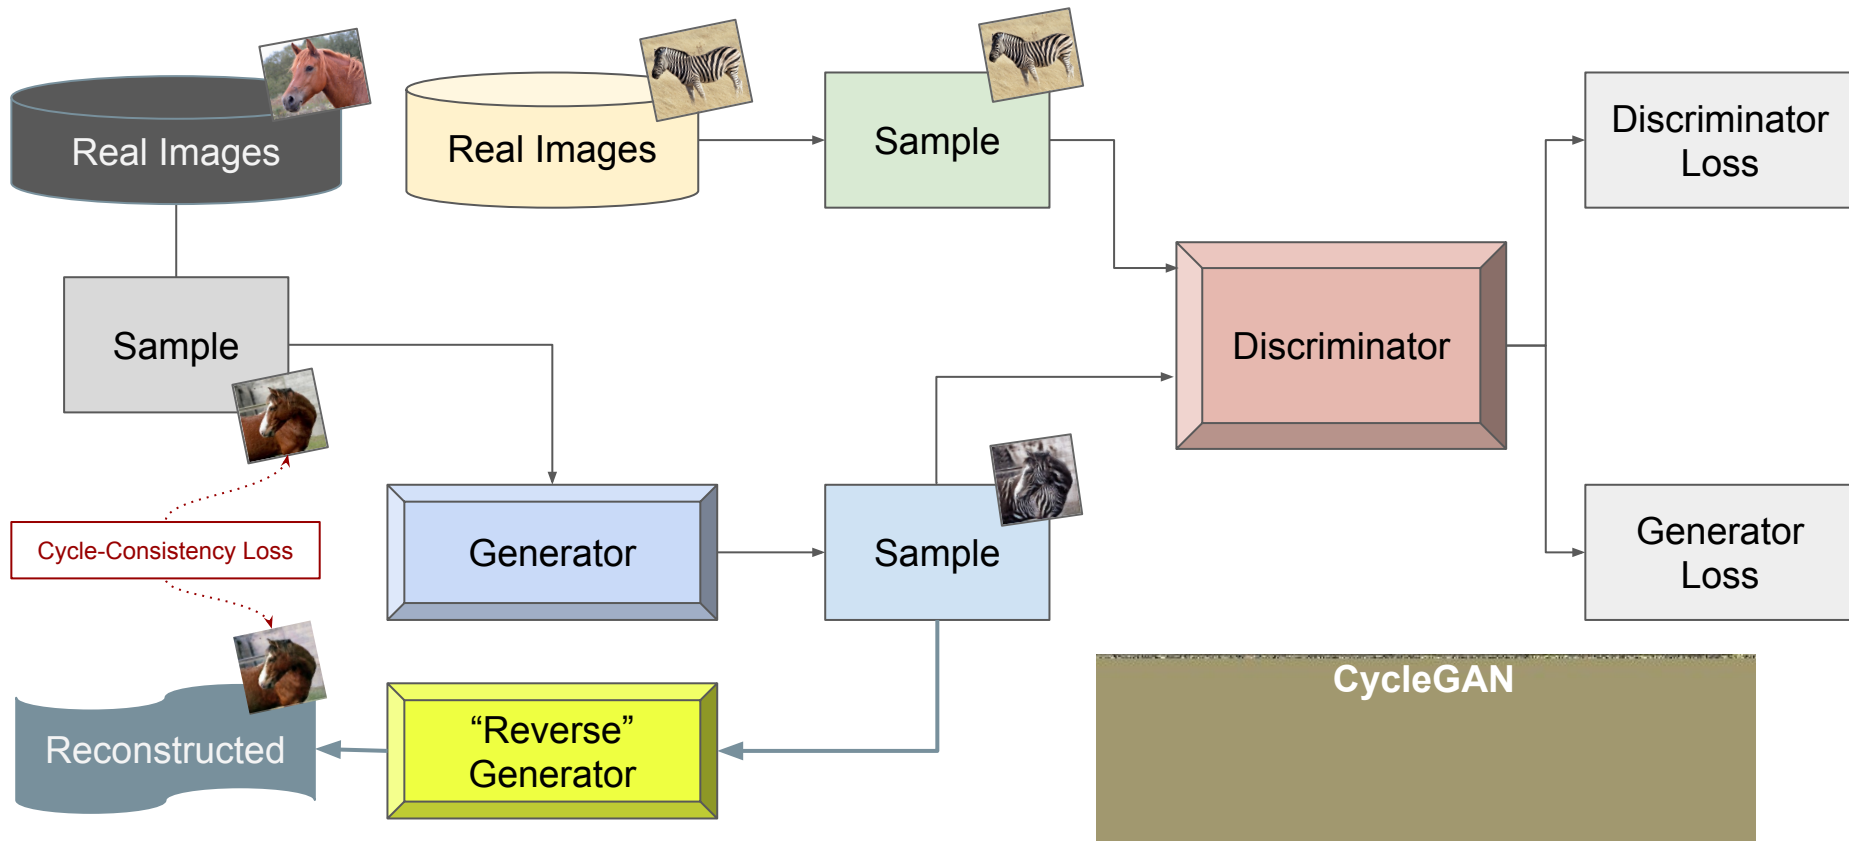
SAGAN: DCGANs + Self-Attention Layer [\(Zhang, et al. 2018\)](#)



Conditional GAN: generate images with specific class



CycleGAN: *unsupervised* conditional GAN



Some GAN loss function variations

SGAN (non-saturating)

$$L_D^{SGAN} = -\mathbb{E}_{x_r \sim \mathbb{P}} [\log(\text{sigmoid}(C(x_r)))] - \mathbb{E}_{x_f \sim \mathbb{Q}} [\log(1 - \text{sigmoid}(C(x_f)))]$$

$$L_G^{SGAN} = -\mathbb{E}_{x_f \sim \mathbb{Q}} [\log(\text{sigmoid}(C(x_f)))]$$

RSGAN

$$L_D^{RSGAN} = -\mathbb{E}_{(x_r, x_f) \sim (\mathbb{P}, \mathbb{Q})} [\log(\text{sigmoid}(C(x_r) - C(x_f)))]$$

$$L_G^{RSGAN} = -\mathbb{E}_{(x_r, x_f) \sim (\mathbb{P}, \mathbb{Q})} [\log(\text{sigmoid}(C(x_f) - C(x_r)))]$$

RaSGAN

$$L_D^{RaSGAN} = -\mathbb{E}_{x_r \sim \mathbb{P}} [\log(\tilde{D}(x_r))] - \mathbb{E}_{x_f \sim \mathbb{Q}} [\log(1 - \tilde{D}(x_f))]$$

$$L_G^{RaSGAN} = -\mathbb{E}_{x_f \sim \mathbb{Q}} [\log(\tilde{D}(x_f))] - \mathbb{E}_{x_r \sim \mathbb{P}} [\log(1 - \tilde{D}(x_r))]$$

$$\tilde{D}(x_r) = \text{sigmoid}(C(x_r) - \mathbb{E}_{x_f \sim \mathbb{Q}} C(x_f))$$

$$\tilde{D}(x_f) = \text{sigmoid}(C(x_f) - \mathbb{E}_{x_r \sim \mathbb{P}} C(x_r))$$

LSGAN

$$L_D^{LSGAN} = \mathbb{E}_{x_r \sim \mathbb{P}} [(C(x_r) - 0)^2] + \mathbb{E}_{x_f \sim \mathbb{Q}} [(C(x_f) - 1)^2]$$

$$L_G^{LSGAN} = \mathbb{E}_{x_f \sim \mathbb{Q}} [(C(x_f) - 0)^2]$$

RaLSGAN

$$L_D^{RaLSGAN} = \mathbb{E}_{x_r \sim \mathbb{P}} [(C(x_r) - \mathbb{E}_{x_f \sim \mathbb{Q}} C(x_f) - 1)^2] + \mathbb{E}_{x_f \sim \mathbb{Q}} [(C(x_f) - \mathbb{E}_{x_r \sim \mathbb{P}} C(x_r) + 1)^2]$$

$$L_G^{RaLSGAN} = \mathbb{E}_{x_f \sim \mathbb{P}} [(C(x_f) - \mathbb{E}_{x_r \sim \mathbb{P}} C(x_r) - 1)^2] + \mathbb{E}_{x_r \sim \mathbb{P}} [(C(x_r) - \mathbb{E}_{x_f \sim \mathbb{Q}} C(x_f) + 1)^2]$$

HingeGAN

$$L_D^{HingeGAN} = \mathbb{E}_{x_r \sim \mathbb{P}} [\max(0, 1 - C(x_r))] + \mathbb{E}_{x_f \sim \mathbb{Q}} [\max(0, 1 + C(x_f))]$$

$$L_G^{HingeGAN} = -\mathbb{E}_{x_f \sim \mathbb{Q}} [C(x_f)]$$

RaHingeGAN

$$L_D^{HingeGAN} = \mathbb{E}_{x_r \sim \mathbb{P}} [\max(0, 1 - \tilde{D}(x_r))] + \mathbb{E}_{x_f \sim \mathbb{Q}} [\max(0, 1 + \tilde{D}(x_f))]$$

$$L_G^{HingeGAN} = \mathbb{E}_{x_f \sim \mathbb{P}} [\max(0, 1 - \tilde{D}(x_f))] + \mathbb{E}_{x_r \sim \mathbb{Q}} [\max(0, 1 + \tilde{D}(x_r))]$$

$$\tilde{D}(x_r) = C(x_r) - \mathbb{E}_{x_f \sim \mathbb{Q}} C(x_f)$$

$$\tilde{D}(x_f) = C(x_f) - \mathbb{E}_{x_r \sim \mathbb{P}} C(x_r)$$

WGAN-GP

$$L_D^{WGAN-GP} = -\mathbb{E}_{x_r \sim \mathbb{P}} [C(x_r)] + \mathbb{E}_{x_f \sim \mathbb{Q}} [C(x_f)] + \lambda \mathbb{E}_{\hat{x} \sim \mathbb{P}_{\hat{x}}} [(\|\nabla_{\hat{x}} C(\hat{x})\|_2 - 1)^2]$$

$$L_G^{WGAN-GP} = -\mathbb{E}_{x_f \sim \mathbb{Q}} [C(x_f)]$$

$\mathbb{P}_{\hat{x}}$ is the distribution of $\hat{x} = \epsilon x_r + (1 - \epsilon)x_f$, where $x_r \sim \mathbb{P}$, $x_f \sim \mathbb{Q}$, $\epsilon \sim U[0, 1]$.

RSGAN-GP

$$L_D^{RSGAN} = -\mathbb{E}_{(x_r, x_f) \sim (\mathbb{P}, \mathbb{Q})} [\log(\text{sigmoid}(C(x_r) - C(x_f)))] + \lambda \mathbb{E}_{\hat{x} \sim \mathbb{P}_{\hat{x}}} [(\|\nabla_{\hat{x}} C(\hat{x})\|_2 - 1)^2]$$

$$L_G^{RSGAN} = -\mathbb{E}_{(x_r, x_f) \sim (\mathbb{P}, \mathbb{Q})} [\log(\text{sigmoid}(C(x_f) - C(x_r)))]$$

$\mathbb{P}_{\hat{x}}$ is the distribution of $\hat{x} = \epsilon x_r + (1 - \epsilon)x_f$, where $x_r \sim \mathbb{P}$, $x_f \sim \mathbb{Q}$, $\epsilon \sim U[0, 1]$.

RaSGAN-GP

$$L_D^{RaSGAN} = -\mathbb{E}_{x_r \sim \mathbb{P}} [\log(\tilde{D}(x_r))] - \mathbb{E}_{x_f \sim \mathbb{Q}} [\log(1 - \tilde{D}(x_f))] + \lambda \mathbb{E}_{\hat{x} \sim \mathbb{P}_{\hat{x}}} [(\|\nabla_{\hat{x}} C(\hat{x})\|_2 - 1)^2]$$

$$L_G^{RaSGAN} = -\mathbb{E}_{x_f \sim \mathbb{Q}} [\log(\tilde{D}(x_f))] - \mathbb{E}_{x_r \sim \mathbb{P}} [\log(1 - \tilde{D}(x_r))]$$

$$\tilde{D}(x_r) = \text{sigmoid}(C(x_r) - \mathbb{E}_{x_f \sim \mathbb{Q}} C(x_f))$$


$$\tilde{D}(x_f) = \text{sigmoid}(C(x_f) - \mathbb{E}_{x_r \sim \mathbb{P}} C(x_r))$$

$\mathbb{P}_{\hat{x}}$ is the distribution of $\hat{x} = \epsilon x_r + (1 - \epsilon)x_f$, where $x_r \sim \mathbb{P}$, $x_f \sim \mathbb{Q}$, $\epsilon \sim U[0, 1]$.

Challenges in GANs

- Mode collapse

Generator produces samples with a limited set of modes




[Source](#)

- ✓ Wasserstein loss
- ✓ Unrolled and packing

- Convergence failure

- ✓ Adding noise to discriminator inputs
- ✓ Penalizing discriminator weights
- ✓ Relativistic metrics

- Vanishing gradient



- ✓ Gradient Penalty
- ✓ Spectral Normalization

- Result evaluation

- ✓ Inception Score
- ✓ Fréchet Inception Distance (FID, MiFID)

Survey

bit.ly/survey_gan